FinOps is not a dashboard problem.

It is a decision-latency problem.
This deck outlines the architecture and operating rhythm that closes the loop in 90 days.

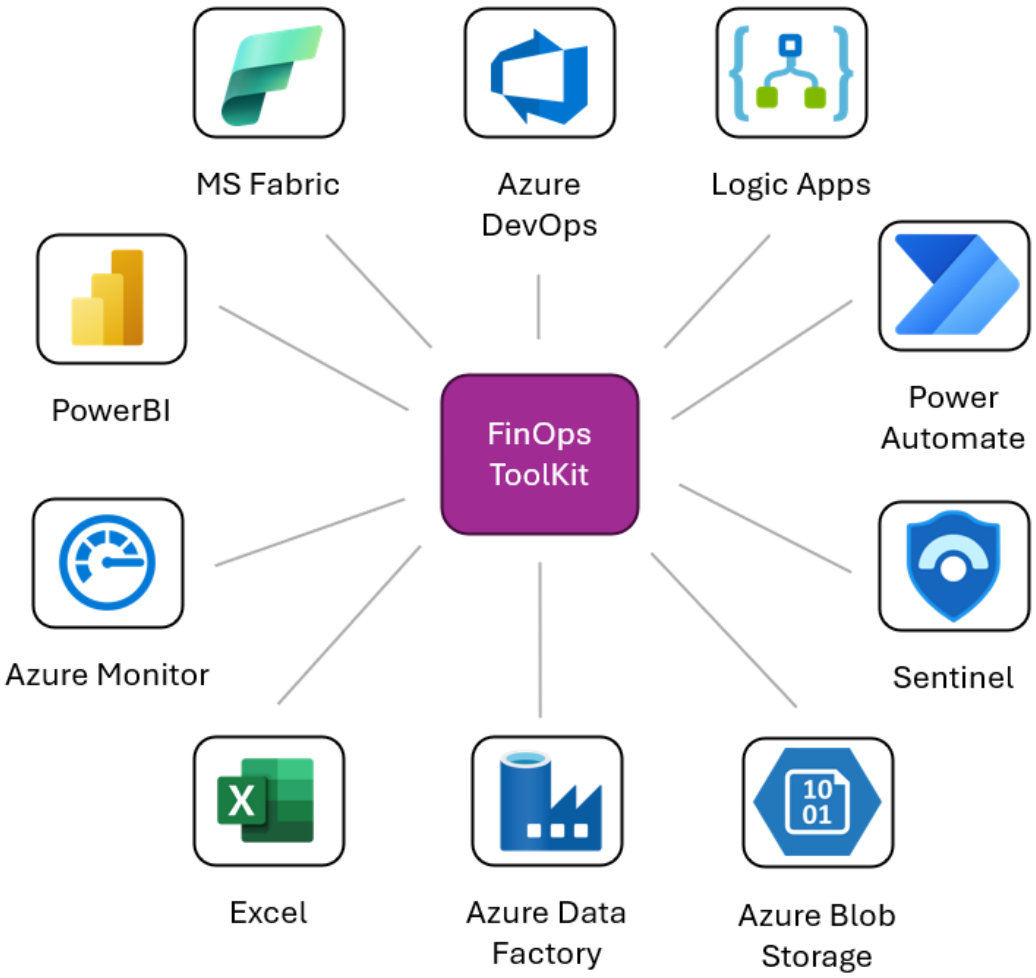

Full article https://forcewake.me/finops-toolkit-framework-playbook-secure-hubs-ai-agents-and-a-90-day-execution-model/

Three non-obvious truths from real FinOps programs

01

Trust data contracts before building dashboards

Most failures come from export scope overlap, missing prices, and ingestion blind spots, not from visualization tooling.

02

Treat FinOps hubs as a platform, not a report pack

The value comes from a stable ingestion model and query layer that supports repeated decisions, not ad hoc analysis.

03

Optimization needs an operating cadence

Recommendations create value only when ownership, thresholds, and change execution are explicit.

Choose architecture by decision latency and scale

Option	Use when	Constraint	Best fit
Raw exports + Power BI	Need visibility this week	Outgrows under complex questions	Small orgs / fast start
Hubs (storage only)	Need standardized baseline	Limited advanced analytics	Centralized reporting
Hubs + ADX	Need KQL speed at enterprise scale	Requires ADX operations capability	Large estates
Hubs + Fabric RTI	Already invested in Fabric model	Capacity + governance overhead	Unified analytics stack

Practical threshold If you rely on KQL-driven reporting and rising cloud spend, move from storage-only exports to a query engine early.

FinOps hubs

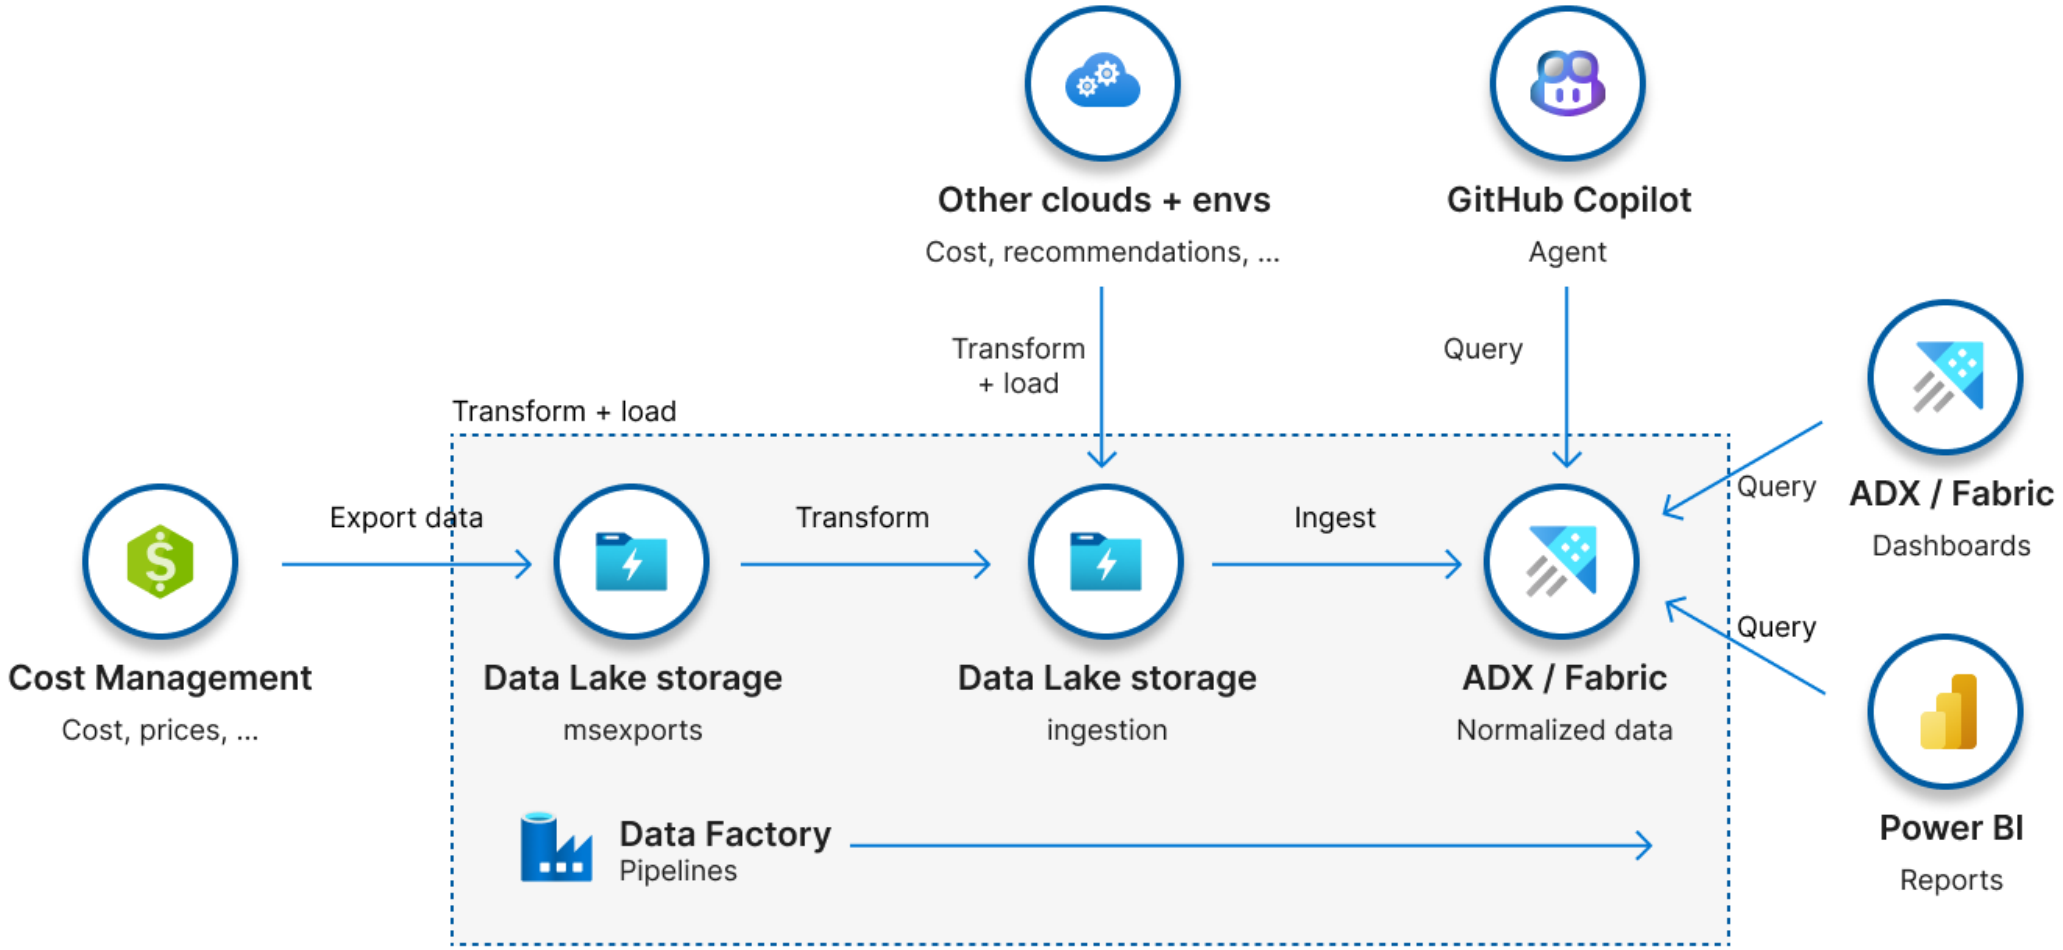
Other clouds + envs
Cost, recommendations, ...
GitHub Copilot
Agent
Transform
+ load
Query
Transform + load
Export data
Transform
Ingest
Query
Query
Cost Management
Cost, prices, ...
Data Lake storage
msexports
Data Lake storage
ingestion
ADX / Fabric
Normalized data
ADX / Fabric
Dashboards
Power BI
Reports
Data Factory
Pipelines

Export -> Ingest -> Model -> Analyze -> Act

Analysis lane: Power BI and workbooks drive weekly cost decisions

AI lane: MCP clients accelerate question-to-insight with read-only defaults

Action lane: FinOps alerts and optimization workflows execute approved changes

Where implementations usually fail and how to prevent it

Failure patterns

- Overlapping export scopes duplicate cost data
- Private networking enabled without DNS readiness
- PriceSheet exported from incorrect MCA scope
- Backfill order reversed (cost before price)

Control patterns

- Version export config and dataset contracts
- Enforce mutually exclusive scopes by design
- Monitor msexports + trigger health as SLO
- Run duplicate checks before each rollout stage

Two commands every implementation team should operationalize

Deploy hub as code

```
Deploy-FinOpsHub -Name "finops-hub-prod" -ResourceGroup "rg-finops" -Location "westeurope"
```

Create daily FocusCost export

```
New-FinOpsCostExport -Name "focuscost-daily" -Scope $scope -StorageAccountId $storageAccountId
-StorageContainer "msexports" -StoragePath "focuscost/<scope-id>" -Dataset FocusCost
-DatasetVersion "1.2-preview" -DoNotOverwrite -Execute
```

Define ownership and thresholds before automation

Metric	Target	Owner	Computation
Data freshness lag	< 24h	Platform Ops	max(x_IngestionTime)
Allocation coverage	> 85%	App teams	Allocated / Total EffectiveCost
Effective savings rate	Trend up	FinOps + engineering	(ListCost - EffectiveCost) / ListCost
Commitment utilization	> 90%	Central FinOps	ConsumedQuantity / PurchasedQuantity

Design rule: promote KQL templates to version-controlled assets and review KPI logic with the same rigor as application code.

Ship in gates: trust -> ownership -> action

Days 1-14

Trust the platform

- Deploy hub + first prod scope
- Validate freshness and schema

Days 15-45

Make ownership real

- Enforce tag contract
- Run weekly anomaly + allocation reviews

Days 46-90

Move to action

- Enable alerts + AOE in read-only mode
- Track recommendation closure rates

Implementation is the moat.

If you want the full playbook with references, code examples, and source mapping:

https://forcewake.me/finops-toolkit-framework-playbook-secure-hubs-ai-agents-and-a-90-day-execution-model/

Next step

- Run a current-state audit against these control points
- Select your target data platform and migration path
- Define KPI owners before enabling optimization automation